

```
In [56]: import pandas as pd
```

```
In [3]: import matplotlib.pyplot as plt
```

```
In [4]: import seaborn as sns
```

```
In [6]: data=pd.read_csv("D:/Riya/SIMPLILEARN/DS Capstone/My project Healthcare/Project_2/health care diabetes.csv")
data.head()
```

Out[6]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

## Week1

```
In [7]: data.isnull().any()
```

```
Out[7]: Pregnancies      False
Glucose      False
BloodPressure  False
SkinThickness  False
Insulin      False
BMI          False
DiabetesPedigreeFunction  False
Age          False
Outcome      False
dtype: bool
```

```
In [8]: data.describe().transpose()
```

```
Out[8]:
```

	count	mean	std	min	25%	50%	75%	max
<b>Pregnancies</b>	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	17.00
<b>Glucose</b>	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.25000	199.00
<b>BloodPressure</b>	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80.00000	122.00
<b>SkinThickness</b>	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32.00000	99.00
<b>Insulin</b>	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127.25000	846.00
<b>BMI</b>	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36.60000	67.10
<b>DiabetesPedigreeFunction</b>	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	2.42
<b>Age</b>	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	81.00
<b>Outcome</b>	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	1.00

```
In [9]: positive=data[data['Outcome']==1]
positive.head()
```

```
Out[9]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>0</b>	6	148	72	35	0	33.6	0.627	50	1
<b>2</b>	8	183	64	0	0	23.3	0.672	32	1
<b>4</b>	0	137	40	35	168	43.1	2.288	33	1
<b>6</b>	3	78	50	32	88	31.0	0.248	26	1
<b>8</b>	2	197	70	45	543	30.5	0.158	53	1

```
In [10]: data['Glucose'].value_counts().head(10)
```

```
Out[10]: 100    17
          99     17
          129    14
          125    14
          111    14
          106    14
           95    13
          108    13
          105    13
          102    13
          Name: Glucose, dtype: int64
```


```
In [11]: plt.hist(data['Glucose'],color='Brown')
```

```
Out[11]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
          array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
                179.1, 199. ]),
          <BarContainer object of 10 artists>)
```




```
In [12]: plt.hist(data['BloodPressure'],color='Green')
```

```
Out[12]: (array([ 35.,   1.,   2.,  13., 107., 261., 243.,  87.,  14.,   5.]),  
array([  0.,  12.2,  24.4,  36.6,  48.8,  61. ,  73.2,  85.4,  97.6,  
        109.8, 122. ]),  
<BarContainer object of 10 artists>)
```




```
In [13]: plt.hist(data['SkinThickness'],color='DarkRed')
```

```
Out[13]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),  
array([ 0. , 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),  
<BarContainer object of 10 artists>)
```




```
In [14]: plt.hist(data['Insulin'],color='Green')
```

```
Out[14]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),  
array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,  
       761.4, 846. ]),  
<BarContainer object of 10 artists>)
```



```
In [15]: plt.hist(data['BMI'],color='Blue')
```

```
Out[15]: (array([ 11.,   0.,  15., 156., 268., 224.,  78.,  12.,   3.,   1.]),  
          array([ 0.   ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,  
                60.39, 67.1 ]),  
          <BarContainer object of 10 artists>)
```



```
In [16]: print(data['Glucose'].value_counts().head())
print(data['BloodPressure'].value_counts().head())
print(data['SkinThickness'].value_counts().head())
print(data['Insulin'].value_counts().head())
data['BMI'].value_counts().head()
```

```
100    17
99     17
129    14
125    14
111    14
Name: Glucose, dtype: int64
70     57
74     52
68     45
78     45
72     44
Name: BloodPressure, dtype: int64
0      227
32     31
30     27
27     23
23     22
Name: SkinThickness, dtype: int64
0      374
105     11
140      9
130      9
120      8
Name: Insulin, dtype: int64
```

```
Out[16]: 32.0    13
31.6     12
31.2     12
0.0      11
33.3     10
Name: BMI, dtype: int64
```


## Week2



```
In [17]: print(positive['Glucose'].value_counts().head())  
plt.hist(positive['Glucose'],histtype='stepfilled',bins=20)
```

```
125    7  
158    6  
128    6  
115    6  
129    6  
Name: Glucose, dtype: int64
```


```
Out[17]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,  
                27., 29., 22., 24., 21., 25., 14.]),  
array([ 0. ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,  
        79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,  
        159.2 , 169.15, 179.1 , 189.05, 199.  ]),  
[<matplotlib.patches.Polygon at 0x24ecb5ff070>])
```



```
In [18]: print(positive['BloodPressure'].value_counts().head())  
plt.hist(positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
70    23  
76    18  
78    17  
74    17  
72    16  
Name: BloodPressure, dtype: int64
```


```
Out[18]: (array([16.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  6.,  6., 19., 37., 56.,  
                36., 41., 31.,  7.,  4.,  4.,  3.]),  
array([ 0. ,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,  
        51.3, 57. , 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,  
        102.6, 108.3, 114. ]),  
[<matplotlib.patches.Polygon at 0x24ecb657d30>])
```



```
In [19]: print(positive['SkinThickness'].value_counts().head())  
plt.hist(positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
0      88  
32     14  
33      9  
30      9  
39      8  
Name: SkinThickness, dtype: int64
```

```
Out[19]: (array([88.,  1.,  4., 10., 18., 30., 41., 34., 23., 15.,  1.,  1.,  1.,  
                0.,  0.,  0.,  0.,  0.,  0.,  1.]),  
array([ 0.   ,  4.95,  9.9 , 14.85, 19.8 , 24.75, 29.7 , 34.65, 39.6 ,  
        44.55, 49.5 , 54.45, 59.4 , 64.35, 69.3 , 74.25, 79.2 , 84.15,  
        89.1 , 94.05, 99.  ]),  
[<matplotlib.patches.Polygon at 0x24ecb6a9f40>])
```



```
In [20]: print(positive['Insulin'].value_counts().head())  
plt.hist(positive['Insulin'],histtype='stepfilled',bins=20)
```

```
0      138
```

```
130      6
```


```
180      4
```

```
156      3
```

```
175      3
```

```
Name: Insulin, dtype: int64
```


```
Out[20]: (array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,  
                5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),  
array([ 0., 42.3, 84.6, 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,  
       380.7, 423., 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,  
       761.4, 803.7, 846. ]),  
[<matplotlib.patches.Polygon at 0x24ecb6feeb0>])
```



```
In [21]: print(positive['BMI'].value_counts().head())  
plt.hist(positive['BMI'],histtype='stepfilled',bins=20)
```

```
32.9    8  
31.6    7  
33.3    6  
30.5    5  
32.0    5  
Name: BMI, dtype: int64
```


```
Out[21]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,  
                14.,  7.,  3.,  1.,  1.,  0.,  1.]),  
         array([ 0.    ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,  
                26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,  
                53.68 , 57.035, 60.39 , 63.745, 67.1  ]),  
         [<matplotlib.patches.Polygon at 0x24ecb75c400>])
```



```
In [22]: BloodPressure = positive['BloodPressure']  
Glucose = positive['Glucose']  
SkinThickness = positive['SkinThickness']  
Insulin = positive['Insulin']  
BMI = positive['BMI']
```


```
In [23]: sns.scatterplot(x='BloodPressure',y='Insulin',hue='Outcome',data=data)
```

```
Out[23]: <AxesSubplot:xlabel='BloodPressure', ylabel='Insulin'>
```




```
In [24]: sns.scatterplot(x='Glucose',y='Insulin',hue='Outcome',data=data,palette='magma')
```

```
Out[24]: <AxesSubplot:xlabel='Glucose', ylabel='Insulin'>
```



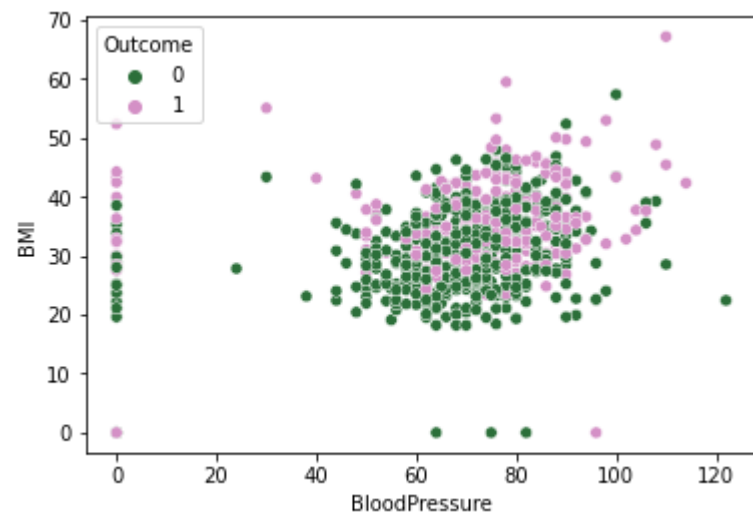
```
In [25]: sns.scatterplot(x='Glucose',y='SkinThickness',hue='Outcome',data=data,palette='rocket_r')
```

```
Out[25]: <AxesSubplot:xlabel='Glucose', ylabel='SkinThickness'>
```




```
In [26]: sns.scatterplot(x='BloodPressure',y='BMI',hue='Outcome',data=data,palette='cubehelix')
```

```
Out[26]: <AxesSubplot:xlabel='BloodPressure', ylabel='BMI'>
```



```
In [27]: sns.scatterplot(x='Insulin',y='BMI',hue='Outcome',data=data,palette='tab10')
```

```
Out[27]: <AxesSubplot:xlabel='Insulin', ylabel='BMI'>
```



```
In [28]: data.corr()
```


```
Out[28]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
<b>Pregnancies</b>	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341
<b>Glucose</b>	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514
<b>BloodPressure</b>	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528
<b>SkinThickness</b>	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970
<b>Insulin</b>	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163
<b>BMI</b>	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242
<b>DiabetesPedigreeFunction</b>	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561
<b>Age</b>	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000
<b>Outcome</b>	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0.238356




```
In [29]: sns.heatmap(data.corr())
```

```
Out[29]: <AxesSubplot:>
```



```
In [30]: plt.subplots(figsize=(8,8))  
sns.heatmap(data.corr(),annot=True,cmap='viridis')
```


Out[30]: <AxesSubplot:>





```
In [31]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True)
```

Out[31]: <AxesSubplot:>



**Week3**

In [32]: `data.head(5)`

Out[32]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

In [33]: `features = data.iloc[:,[0,1,2,3,4,5,6,7]].values`  
`label = data.iloc[:,8].values`

In [34]: `from sklearn.model_selection import train_test_split`  
`X_train,X_test,y_train,y_test = train_test_split(features,`  
`label,`  
`test_size=0.2,`  
`random_state =10)`

In [35]: `from sklearn.linear_model import LogisticRegression`  
`model = LogisticRegression()`  
`model.fit(X_train,y_train)`

Out[35]: `LogisticRegression()`

In [36]: `print(model.score(X_train,y_train))`  
`print(model.score(X_test,y_test))`

0.7719869706840391  
0.7662337662337663

In [37]: `from sklearn.metrics import confusion_matrix`  
`cm = confusion_matrix(label,model.predict(features))`  
`cm`

Out[37]: `array([[446, 54],`  
`[122, 146]], dtype=int64)`

```
In [38]: from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```

	precision	recall	f1-score	support
0	0.79	0.89	0.84	500
1	0.73	0.54	0.62	268
accuracy			0.77	768
macro avg	0.76	0.72	0.73	768
weighted avg	0.77	0.77	0.76	768


In [39]: *#Preparing ROC Curve (Receiver Operating Characteristics Curve)*

```
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.837

Out[39]: [





```
In [40]: #Applying Decision Tree Classifier  
from sklearn.tree import DecisionTreeClassifier  
model3 = DecisionTreeClassifier(max_depth=5)  
model3.fit(X_train,y_train)
```

Out[40]: DecisionTreeClassifier(max\_depth=5)

```
In [41]: model3.score(X_train,y_train)
```

Out[41]: 0.8289902280130294

```
In [42]: model3.score(X_test,y_test)
```

Out[42]: 0.7727272727272727

```
In [43]: #Applying Random Forest  
from sklearn.ensemble import RandomForestClassifier  
model4 = RandomForestClassifier(n_estimators=11)  
model4.fit(X_train,y_train)
```

Out[43]: RandomForestClassifier(n\_estimators=11)

```
In [44]: model4.score(X_train,y_train)
```

Out[44]: 0.99185667752443

```
In [45]: model4.score(X_test,y_test)
```

Out[45]: 0.7532467532467533

```
In [46]: #Support Vector Classifier  
  
from sklearn.svm import SVC  
model5 = SVC(kernel='rbf',  
              gamma='auto')  
model5.fit(X_train,y_train)
```

Out[46]: SVC(gamma='auto')

```
In [49]: #Applying K-NN  
from sklearn.neighbors import KNeighborsClassifier  
model2 = KNeighborsClassifier(n_neighbors=7, metric='minkowski', p = 2)  
model2.fit(X_train,y_train)
```

```
Out[49]: KNeighborsClassifier(n_neighbors=7)
```

```

In [50]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".format(tpr, fpr, thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")

```


AUC: 0.836

```

True Positive Rate - [0.      0.06716418 0.23880597 0.44776119 0.60074627 0.75373134
 0.88059701 0.98507463 1.      ], False Positive Rate - [0.      0.      0.02  0.056 0.12  0.248 0.428 0.668 1.
] Thresholds - [2.      1.      0.85714286 0.71428571 0.57142857 0.42857143
 0.28571429 0.14285714 0.      ]

```

Out[50]: Text(0, 0.5, 'True Positive Rate')




```
In [51]: #Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.624 auc=0.726 ap=0.727

Out[51]: [<matplotlib.lines.Line2D at 0x24ecda306a0>]




In [52]: *#Precision Recall Curve for KNN*

```
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.658 auc=0.752 ap=0.709

Out[52]: [`<matplotlib.lines.Line2D at 0x24ecda8f7c0>`]






In [53]: *#Precision Recall Curve for Decision Tree Classifier*

```
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model3.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model3.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.714 auc=0.815 ap=0.768

Out[53]: [




```
In [54]: #Precision Recall Curve for Random Forest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model4.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.917 auc=0.962 ap=0.954

```
Out[54]: [<matplotlib.lines.Line2D at 0x24ecdb44a60>]
```



In [ ]:

In [ ]:

In [ ]:

In [ ]: